

Classification Problem Formulation

Dr. Kalidas Y., IIT Tirupati

By the end of this lecture, you will understand what is a classification problem and be able to formulate simple ones

Loss function

$$L(w) = -\log P(D|w) = - \sum_{(x,y) \in D} (y * (w \cdot x) - \log(1 + e^{w \cdot x}))$$

44) key phrase... “Classification Problem”

- Category
- Examples
 - Cat or Not Cat
 - Good or Bad
 - Rainy Day, Cloudy Day, Bright Day, Warm Day, Hot Day
 - Tall, Moderately Tall, Short
 - Etc.

45) key phrase... “Binary Classification”

- Two Categories
- Examples
 - Cat or Not Cat
 - Good or Bad

46) key phrase... “Multi-class Classification”

- Multiple Categories
- Examples
 - Fault categories – Valve fault, Temperature fault, Pressure fault
 - Identify of a person – you may have thousands of people and millions of photos
 - Common objects – Car, Bus, Cycle, Person, Animal etc.

47) key phrase... “Threshold”

- Score
- Threshold
- Above it → One class
- Below it → Another class
- *Craft loss function accordingly*

48) key phrase... “SoftMax”

- K scores (*assume some numbers smoothed out in some way for now*)
- Pick maximal score – for e.g. i^{th} is maximum
- Select i^{th} class
- *Craft loss function accordingly*

Example 1 – Logistic Loss Function

- Knowledge of x_i (some k dimensional vector)
 - Knowledge of $y_i \in \{-1, 1\}$
 - Model... $y = f(x) = w \cdot x$
 - Loss function, $L(w) = \sum_{i=1}^N \log(1 + e^{-y_i \times w \cdot x_i})$
 - Data set, $\{(x_1, y_1), \dots, (x_N, y_N)\}$
-
- What happens when y_i and $w \cdot x_i$ have same sign?
 - What happens when y_i and $w \cdot x_i$ have opposite sign?

Example 2 – Hinge Loss Function

- Knowledge of x_i (some k dimensional vector)
- Knowledge of $y_i \in \{-1, 1\}$
- Model... $y = f(x) = w \cdot x$
- Loss function, $L(w) = \frac{\|w\|_2^2}{2} + \sum_{i=1}^N \max(0, 1 - y_i \times w \cdot x_i)$
- Data set, $\{(x_1, y_1), \dots, (x_N, y_N)\}$

- What happens when y_i and $w \cdot x_i$ have same sign?
- What happens when y_i and $w \cdot x_i$ have opposite sign?

$$E(w) = \frac{1}{N} \sum_{i=1}^{i=N} L(y_i, f(x_i)) + \alpha R(w)$$

Annotations:

- $E(w)$: Regularized loss function
- $\frac{1}{N}$: general practice [optional] to avoid numerical precision issues
- $\sum_{i=1}^{i=N}$: Loss function
- y_i : ground truth
- $f(x_i)$: model
- α : control amount of regularization
- $R(w)$: regularization e.g. L1, L2

Popular Loss Functions

1. Hinge Loss: $L(y_i, f(x_i)) = \max(0, 1 - y_i \times f(x_i))$
 2. Perceptron Loss: $L(y_i, f(x_i)) = \max(0, -y_i \times f(x_i))$
 3. Huber Loss: $L(y_i, f(x_i)) = \epsilon |y_i - f(x_i)| - \frac{1}{2} \epsilon^2$
 4. Modified Huber Loss: $L(y_i, f(x_i)) = \max(0, 1 - y_i \times f(x_i))^2$
 5. Logistic Loss: $L(y_i, f(x_i)) = \log(1 + e^{-y_i \times f(x_i)})$
 6. Least Squares Loss: $L(y_i, f(x_i)) = \frac{1}{2} (y_i - f(x_i))^2$
 7. Epsilon-Insensitive Loss: $L(y_i, f(x_i)) = \max(0, |y_i - f(x_i)| - \epsilon)$
- Regularization term
 1. Lasso Regularization: $R(w) = |w|_1$
 2. Ridge Regularization: $R(w) = |w|_2^2$
 3. Elastic Net Regularization: $R(w) = \beta |w|_2^2 + (1 - \beta) |w|_1$

REF - https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html

```
In [43]: import numpy as np
        from sklearn.linear_model import SGDClassifier
```

```
In [44]: X = np.array([[-1, -1], [-2, -1], [1, 1], [2, 1]])
        Y = np.array([1, 1, 2, 2])
```

```
In [45]: clf = SGDClassifier(loss='log', fit_intercept=False) #try other values 'hinge', 'squared_loss'
```

```
In [46]: clf.fit(X,Y)
```

```
Out[46]: SGDClassifier(alpha=0.0001, average=False, class_weight=None, epsilon=0.1,
                        eta0=0.0, fit_intercept=False, l1_ratio=0.15,
                        learning_rate='optimal', loss='log', max_iter=5, n_iter=None,
                        n_jobs=1, penalty='l2', power_t=0.5, random_state=None,
                        shuffle=True, tol=None, verbose=0, warm_start=False)
```

```
In [47]: print(clf.predict([[-0.8, -1]]))

[1]
```

```
In [48]: X2 = X + 100
```

```
In [49]: clf.fit(X2,Y)
```

```
Out[49]: SGDClassifier(alpha=0.0001, average=False, class_weight=None, epsilon=0.1,
                        eta0=0.0, fit_intercept=False, l1_ratio=0.15,
                        learning_rate='optimal', loss='log', max_iter=5, n_iter=None,
                        n_jobs=1, penalty='l2', power_t=0.5, random_state=None,
                        shuffle=True, tol=None, verbose=0, warm_start=False)
```

```
In [50]: X_test = np.array([[-0.8,1]]) + 100
        print(clf.predict(X_test))

[1]
```

When to choose what loss function?

It is a hyper parameter as well...

We will see a bit more about these in the subsequent lecture!

49) key phrase... “hyper parameters & search”

- Each and every parameter that a human can adjust
- Those parameters which cannot be learned from data automatically
- Those knobs that require human-in-loop way of adjustment
- Example
 - 1) Degree of a polynomial in case of polynomial fitting problem
 - 2) Type of regularization to use – L1 or L2 or Elastic Net
 - 3) The multiplication factor for regularization term
 - 4) Type of loss function to use – you have about 7 popular types and even more...
 - 5) In case of tree based methods... number of trees, depth of each tree etc. (will see)

*For all those knobs in your hand... try different values.. **a.k.a GRID SEARCH***

50) key phrase... “Grid Search”

- Programatically exploring possible value settings for “hyper parameters”
- to choose the best one
- Best one in terms of the highest scoring
- Or least error
- That setting which gives highest score or least error